

OPEN ENDED LAB REPORT

DESIGN A MODEL TO FORECAST IN THE CITYLEARN CHALLENGE FORECASTING NON-SHIFTABLE LOAD USING MACHINE LEARNING ON THE CITYLEARN DATASET

SUBMITTED BY :

MAHAVIA KHAN

21KTELE0637

UNIVERSITY OF ENGINEERING AND TECHNOLOGY PESHAWAR,
JALOZAI CAMPUS
DEPARTMENT OF ELECTRICAL ENGINEERING
LAB: MACHINE LEARNING
OPENENDED LAB REPORT

SUBMITTED TO :

ENGR.IRSHAD ULLHA


DATE OF SUBMISSION: JUNE, 2025

Contents

1	Introduction	3
2	Objectives	5
3	Data Overview	6
3.1	Building 2 Dataset	6
3.2	Carbon Intensity Dataset	6
3.3	Pricing Dataset	7
3.4	Weather Dataset	7
4	Data Preprocessing	8
5	Model Architecture	10
5.1	Data Preparation and Sequencing	10
5.2	Model Structure	11
5.2.1	Model Layers	11
5.2.2	Model Code	11
5.2.3	Additional Model Variants	12
5.3	Training Setup	13
5.3.1	Loss Function and Optimization	13
5.3.2	Training Configuration	13
5.3.3	Training Procedure	14
6	Evaluation and Results	15
6.1	Evaluation Metrics	15
7	Future Work and Limitations	16
7.1	Future Work	16
7.2	Limitations	16
8	Conclusion	17
9	References	18

INTRODUCTION

The global transition toward smarter, more sustainable energy systems has amplified the importance of accurate load forecasting in electrical grids. As electricity demand becomes increasingly dynamic—driven by the integration of renewable energy sources, electric vehicles, and decentralized energy production—utilities are under greater pressure to ensure efficient, stable, and responsive energy distribution. A critical component in this complex landscape is the ability to predict non-shiftable electrical loads—those essential and inflexible demands that cannot be deferred or rescheduled without compromising user needs or safety.



Non-shiftable loads—such as lighting, heating, refrigeration, and life-support systems—represent a foundational portion of daily energy consumption. These loads are consistent and time-bound, making their accurate prediction vital for real-time supply-demand balancing, peak load mitigation, and optimal allocation of energy resources. Traditional forecasting methods, often based on statistical techniques or historical averaging, typically fall short in capturing the complex, non-linear relationships inherent in modern consumption patterns—especially when influenced by external factors such as weather, occupant behavior, and pricing signals.

With the advancement of machine learning (ML) and deep learning, there is now a transformative opportunity to improve the accuracy and flexibility of load forecasting. ML models, particularly those suited for time-series forecasting, can learn hidden patterns in sequential data and incorporate diverse contextual features to make data-driven, future-aware predictions. Architectures like Long Short-Term Memory (LSTM) networks have shown promising results due to their ability to capture long-range temporal dependencies—making them especially

useful in environments where energy demand changes dynamically over time.

This report explores a machine learning-based approach to forecasting non-shiftable load within the framework of the CityLearn environment—an open-source reinforcement learning platform designed for simulating energy coordination and demand response in smart urban districts. While CityLearn primarily focuses on controlling shiftable loads using reinforcement learning agents, a crucial parallel task is the accurate forecasting of non-shiftable loads, which directly impacts how effectively these agents can perform.

By designing a dedicated ML model to forecast non-shiftable consumption in the CityLearn Challenge, this work addresses a core need in smart energy systems: ensuring essential demand is met reliably and efficiently. The challenge not only requires precise model selection, but also emphasizes preprocessing, feature engineering, and guarding against data leakage. In doing so, it reflects the broader industry shift toward data-driven energy management, contributing to the global goal of building resilient, low-carbon, and intelligent power infrastructures.

OBJECTIVES

The primary objective of this project is to develop an effective and reliable machine learning model capable of forecasting non-shiftable electrical load in a smart grid environment. By leveraging time-series data and contextual features such as weather conditions, electricity pricing, and carbon intensity, the model aims to enhance the predictability of essential, inflexible energy demand across urban buildings. The key objectives of this project are as follows:

- To investigate the characteristics of non-shiftable load and its role in modern energy systems.
- To explore and preprocess the CityLearn dataset, focusing on features relevant to non-shiftable load forecasting.
- To engineer temporal, environmental, and economic features that contribute significantly to the accuracy of load prediction models.
- To design and implement a deep learning-based forecasting model.
- To evaluate the performance of the model using metrics such as Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and Mean Absolute Percentage Error (MAPE).
- To identify the impact of predicted input variables (e.g., forecasted weather or pricing data) on model performance and forecasting robustness.
- To highlight challenges such as data leakage and dependency on the quality of forecasted inputs, and propose mitigation strategies.

These objectives align with the broader goal of integrating artificial intelligence into smart grid operations, improving energy efficiency, and supporting future-ready infrastructure in urban energy systems.

DATA OVERVIEW

The data used in this project is sourced from the CityLearn Challenge 2022 repository (Phase-All dataset) and consists of four distinct datasets: Building 2, Carbon Intensity, Pricing, and Weather. These datasets collectively provide a rich set of temporal, environmental, economic, and operational features essential for accurate forecasting of non-shiftable electrical load in smart buildings.

3.1 BUILDING 2 DATASET

This dataset contains detailed building-specific measurements and contextual features, including:

- Temporal variables: month, hour, day type (e.g., weekday/week-end), daylight savings status
- Indoor environmental conditions: indoor dry bulb temperature, indoor relative humidity
- Comfort-related features: average unmet cooling setpoint difference
- Energy demands: non-shiftable load (target variable), domestic hot water (DHW) demand, cooling demand, heating demand
- Renewable generation: solar generation data

3.2 CARBON INTENSITY DATASET

Contains time-series data for carbon intensity, representing the carbon emissions associated with electricity generation over time. This helps to understand the environmental impact linked to energy consumption.

3.3 PRICING DATASET

Includes electricity pricing data with both actual prices and short-term predictions, capturing economic signals that may influence consumption behavior. Variables include:

- Actual electricity pricing
- Predicted electricity pricing for the next 1, 2, and 3 time steps ahead

3.4 WEATHER DATASET

Contains observed and forecasted weather parameters influencing building energy consumption:

- Observed features: outdoor dry bulb temperature, outdoor relative humidity, diffuse solar irradiance, direct solar irradiance
- Forecasted features: predicted values for the above variables for 1, 2, and 3 time steps ahead

All datasets are time-synchronized to ensure consistency across features and enable the development of robust forecasting models. Preprocessing steps such as missing data imputation, normalization, and temporal feature engineering were applied to enhance model performance.

The integration of these datasets allows the model to capture complex interactions between building operations, environmental conditions, economic factors, and carbon emissions, facilitating more accurate and reliable forecasting of non-shiftable electrical load in urban smart grids.

DATA PREPROCESSING

This section outlines the preprocessing steps applied to prepare the raw datasets for model development. These steps ensured data consistency, minimized noise, and enhanced the predictive potential of input features.

The process began by loading and merging the four key datasets â **Building 2**, **Weather**, **Carbon Intensity**, and **Pricing** â using a common time index. This integration created a unified dataset encompassing temporal, environmental, and economic variables.

To reduce redundancy and focus on relevant predictors, the following non-essential columns were removed:

- indoor_dry_bulb_temperature
- average_unmet_cooling_setpoint_difference
- indoor_relative_humidity

Subsequently, **outliers** were detected and handled across key numerical features to ensure robust model performance. The columns were **reordered** to logically group features, enhancing readability and consistency during modeling.

Missing values were identified and addressed using appropriate imputation strategies to preserve data quality.

To capture the cyclical nature of temporal features such as hour and month, **cyclical encoding** was applied using sine and cosine transformations. This enabled the model to better understand periodic patterns in the data.

A **correlation matrix** was computed to analyze inter-feature relationships and support feature selection by identifying strongly correlated or irrelevant variables.

The final dataset was partitioned into three sets using a 70:20:10 split ratio:

- Training Set: (876, 26)
- Validation Set: (1752, 26)

- Test Set: (876, 26)

Datetime components were extracted where applicable to enhance temporal context.

For feature scaling, two normalization techniques were applied:

- **StandardScaler** was used for weather-related continuous variables.
- **MinMaxScaler** was applied to economic and target variables.

The scaled data was then saved as three separate sets â training, validation, and testing â for use during model training and evaluation. After model predictions, the target variable `non_shiftable_load` was **inverse transformed** to return results to their original scale, ensuring meaningful interpretation.

MODEL ARCHITECTURE

This section outlines the architecture and implementation of the deep learning model developed for forecasting non-shiftable load using the CityLearn dataset.

The modeling process begins by loading the preprocessed training, validation, and test datasets, each containing normalized values. Using a sliding window approach, the time series data is reshaped into sequences of 24 time steps to predict the next step (1-step ahead). This captures temporal patterns crucial for accurate load forecasting.

5.1 DATA PREPARATION AND SEQUENCING

The datasets are loaded from saved CSV files containing normalized data. A sliding window method is used to generate sequences of length 24 time steps, which are used as model inputs. The target for prediction is the immediate next step following each sequence, enabling the model to learn temporal dependencies. listings xcolor

```

1 # Define dataset directory
2 path_dataset = 'G:\\8TH SEMESTER\\MahaviaKhan_ML\\OpenEndedLab_0637\\
   Dataset_CityLearn'
3 # Load training data
4 path_tr = os.path.join(path_dataset, 'CityLearn_train.csv')
5 df_tr = pd.read_csv(path_tr)
6 train_set = df_tr.iloc[:].values
7 # Load validation data
8 path_v = os.path.join(path_dataset, 'CityLearn_validation.csv')
9 df_v = pd.read_csv(path_v)
10 validation_set = df_v.iloc[:].values
11 # Load test data
12 path_te = os.path.join(path_dataset, 'Citylearn_test.csv')
13 df_te = pd.read_csv(path_te)
14 test_set = df_te.iloc[:].values
15 # Load scaler for feature normalization
16 path_scaler = os.path.join(path_dataset, 'Scaler.pkl')
17 scaler = pickle.load(open(path_scaler, 'rb'))
18 # Print dataset shapes
19 train_set.shape, validation_set.shape, test_set.shape

```

Listing 1: Dataset Loading and Scaling from CityLearn CSV Files

5.2 MODEL STRUCTURE

The model architecture is based on a multi-layer Long Short-Term Memory (LSTM) network designed to handle sequential time-series data efficiently.

5.2.1 Model Layers

Layer	Configuration
Input Layer	Sequence length = 24, Feature dimension = 29
LSTM Layers	3 layers, 60 hidden units each, dropout = 0.2
Layer Normalization	Applied after LSTM output
Fully Connected (Dense)	Linear layer with 1 output

Table 1: Summary of model architecture layers

listings xcolor

5.2.2 Model Code

The following Python code defines the LSTM-based neural network architecture used for forecasting non-shiftable load. The model consists of a three-layer LSTM with dropout for regularization, followed by layer normalization and a fully connected layer for output. The weights are initialized using Xavier initialization for improved training stability.

```

1 import torch
2 import torch.nn as nn
3
4 class MOSTLY_CITED(nn.Module):
5     def __init__(self):
6         super().__init__()
7         self.device = ('cuda' if torch.cuda.is_available() else 'cpu')
8         self.lstm = nn.LSTM(29, 60, 3, batch_first=True, dropout=0.2).to(
9             self.device) # Added dropout
10        self.layer_norm = nn.LayerNorm(60).to(self.device) # Added
11        LayerNorm
12        self.fc = nn.Linear(60, 1).to(self.device)
13
14        for name, param in self.lstm.named_parameters():
15            if 'weight_ih' in name or 'weight_hh' in name:
16                torch.nn.init.xavier_normal_(param.data)
17            elif 'bias' in name:
18                nn.init.zeros_(param)
19
20    def forward(self, x):

```

```

19     h0 = torch.zeros(3, x.size(0), 60).to(self.device)
20     c0 = torch.zeros(3, x.size(0), 60).to(self.device)
21     out, _ = self.lstm(x, (h0, c0))
22     out = self.layer_norm(out[:, -1, :]) # Apply LayerNorm
23     out = self.fc(out) # Removed sigmoid activation
24     return out

```

Listing 2: PyTorch LSTM Model Class: MOSTLY_CITED

5.2.3 Additional Model Variants

To compare different recurrent neural network architectures for time series forecasting, two additional models were implemented: a stacked LSTM and a simple RNN-based architecture. These models explore how architectural depth, memory mechanisms, and hidden layer sizes impact the model's ability to capture temporal dependencies in non-shiftable load forecasting.

Model	Layer	Configuration
LSTMModel	Stacked LSTM	2 layers, hidden sizes = 64 and 32, dropout = 0.3
	Fully Connected	Linear(32 $\rightarrow$ 16), ReLU, Linear(16 $\rightarrow$ 1)
RNNModel	Stacked RNN	2 layers, hidden sizes = 8 and 20
	Flatten + Output	Flatten, Linear(20 $\rightarrow$ 1)

Table 2: Architectural Summary of RNN and LSTM Variants

Model Code Implementations: listings xcolor

```

1  class LSTMModel(nn.Module):
2      def __init__(self, time_steps=24, num_features=29):
3          super(LSTMModel, self).__init__()
4          self.lstm1 = nn.LSTM(input_size=num_features, hidden_size=64,
5                               batch_first=True, dropout=0, bidirectional=
6                                   False)
7          self.dropout1 = nn.Dropout(0.3)
8          self.lstm2 = nn.LSTM(input_size=64, hidden_size=32,
9                               batch_first=True, dropout=0, bidirectional=
10                                   False)
11          self.dropout2 = nn.Dropout(0.3)
12          self.fc1 = nn.Linear(32, 16)
13          self.relu = nn.ReLU()
14          self.output_layer = nn.Linear(16, 1)
15
16      def forward(self, x):
17          x, _ = self.lstm1(x)
18          x = self.dropout1(x)
19          x, _ = self.lstm2(x)
20          x = self.dropout2(x)
21          x = x[:, -1, :] # Last time step
22          x = self.fc1(x)
23          x = self.relu(x)
24          return self.output_layer(x)

```

Listing 3: Stacked LSTM Model Definition

listings xcolor

```

1 class RNNModel(nn.Module):
2     def __init__(self, time_steps=24, num_features=29):
3         super(RNNModel, self).__init__()
4         self.rnn1 = nn.RNN(input_size=num_features, hidden_size=8,
5                             batch_first=True)
6         self.rnn2 = nn.RNN(input_size=8, hidden_size=20, batch_first=True)
7         self.flatten = nn.Flatten()
8         self.output_layer = nn.Linear(20, 1)
9
10    def forward(self, x):
11        x, _ = self.rnn1(x)
12        x, _ = self.rnn2(x)
13        x = x[:, -1, :] # Last time step
14        x = self.flatten(x)
15        return self.output_layer(x)

```

Listing 4: Stacked RNN Model Definition

5.3 TRAINING SETUP

5.3.1 Loss Function and Optimization

The model is trained using Mean Squared Error (MSE) loss function optimized with the Adam optimizer. A learning rate scheduler reduces the learning rate if the validation loss plateaus to fine-tune training.

5.3.2 Training Configuration

Component	Setting / Details
Loss Function	Mean Squared Error (MSE)
Optimizer	Adam
Learning Rate	0.001
Scheduler	ReduceLROnPlateau (factor=0.5, patience=10)
Forecasting Type	Univariate, 1-step ahead
Input Shape	(batch size, 24, 29)
Target Variable	Non-shiftable load

Table 3: Training and optimization details

This architecture effectively captures temporal dependencies in energy consumption data, while techniques such as dropout and layer normalization improve model robustness and generalization.

5.3.3 *Training Procedure*

The training process was initialized with a starting epoch of 1 and set to run for 50 epochs. The model was trained using mini-batches with a batch size of 32 to efficiently manage memory and improve training speed.

Data loaders were created for the training, validation, and test datasets to facilitate batch-wise data feeding into the model.

Mean Squared Error (MSE) loss was used as the primary criterion for optimization, while Mean Absolute Error (MAE) was also prepared for evaluation.

A utility was implemented to save the best model based on validation performance during training. Additionally, training and validation loss curves were plotted to monitor the model's learning progress.

The total training time was recorded to understand the computational cost.

EVALUATION AND RESULTS

This section presents the evaluation of the trained model on the test dataset using standard performance metrics. The purpose is to quantify how well the model predicts the target variable compared to the actual observed values.

6.1 EVALUATION METRICS

To assess the model performance, three widely used error metrics were calculated:

- **Mean Squared Error (MSE):** Measures the average squared difference between predicted and actual values. It heavily penalizes large errors, providing a strong indication of overall prediction accuracy.
- **Root Mean Squared Error (RMSE):** The square root of MSE, giving an error measure in the same units as the target variable. It helps to interpret the magnitude of prediction errors.
- **Mean Absolute Percentage Error (MAPE):** Expresses the prediction error as a percentage of the actual values, providing an intuitive measure of relative error.

The evaluation results are summarized in Table 4 below.

Table 4: Model Evaluation Metrics

Metric	Value
Mean Squared Error (MSE)	1,401,531.84
Root Mean Squared Error (RMSE)	1,183.86
Mean Absolute Percentage Error (MAPE)	2.98 %

FUTURE WORK AND LIMITATIONS

7.1 FUTURE WORK

There are several avenues to extend and improve the current model. First, integrating attention mechanisms such as self-attention or Transformer-based encoders may further enhance temporal dependency modeling, especially for long sequences. Additionally, exploring ensemble models or hybrid architectures that combine CNNs or GRUs with LSTM can be promising for improving forecasting accuracy.

Future work may also involve deploying the trained model in a real-time environment using IoT-based systems. This would enable dynamic load forecasting and timely energy distribution in smart grids. Moreover, using more granular datasets (hourly or 15-min intervals) and additional external factors like temperature, weather, and socio-economic variables could significantly improve prediction performance.

7.2 LIMITATIONS

Despite the promising results, the proposed model has certain limitations. Firstly, the model performance heavily depends on the quality and quantity of input features; insufficient or noisy data may lead to poor generalization. Secondly, although the LSTM model captures sequential dependencies well, it still struggles with long-range dependencies when compared to attention-based models.

Another limitation is computational cost. The training time on standard hardware (CPU/GPU) is relatively high, and hyperparameter tuning further increases the training duration. Lastly, the model was trained and evaluated on historical data only, without incorporating real-time changes in grid behavior or unexpected disruptions, which might limit its robustness in production environments.

CONCLUSION

This project focused on forecasting non-shiftable load using machine learning techniques applied to the CityLearn dataset. A deep learning-based approach was developed with a multi-layer LSTM architecture, enhanced by Layer Normalization and Xavier initialization, to effectively capture the temporal dependencies inherent in non-shiftable energy consumption patterns.

The dataset was carefully preprocessed, and the model was trained over 50 epochs using MSE as the primary loss function. Performance evaluation using MSE, RMSE, and MAPE indicated that the proposed model achieved high accuracy in predicting future non-shiftable loads.

The experimental results validate the effectiveness of LSTM-based architectures in handling sequential energy data, particularly for non-flexible, demand-side load components. The structured pipeline—from data loading to model evaluation—ensures the approach is reproducible and scalable.

This work contributes to the growing field of energy demand forecasting by offering a reliable solution for non-shiftable load prediction, which plays a crucial role in planning, resource allocation, and demand-side energy management strategies.

REFERENCES

-
- Vazquez-Canteli, J. R., & Nagy, Z. (2019). *Reinforcement learning for demand response: A review of algorithms and modeling techniques*. *Applied Energy*, 235, 1072â1089. <https://doi.org/10.1016/j.apenergy.2018.11.002>
- CityLearn GitHub Repository. (2023). *CityLearn: An Open AI Gym Environment for the Implementation of Demand Response Using Reinforcement Learning for Building Energy Management*. Available at: <https://github.com/intelligent-environments-lab/CityLearn>
- Hochreiter, S., & Schmidhuber, J. (1997). *Long short-term memory*. *Neural Computation*, 9(8), 1735â1780. <https://doi.org/10.1162/neco.1997.9.8.1735>
- Armstrong, J. S. (1985). *Long-Range Forecasting: From Crystal Ball to Computer*. Wiley-Interscience.
- Taieb, S. B., & Hyndman, R. J. (2014). *A gradient boosting approach to the Kaggle load forecasting competition*. *International Journal of Forecasting*, 30(2), 382â394. <https://doi.org/10.1016/j.ijforecast.2013.07.005>
- Marino, D. L., Amarasinghe, K., & Manic, M. (2016). *Building Energy Load Forecasting using Deep Neural Networks*. 2016 IEEE IECON. <https://doi.org/10.1109/IECON.2016.7793247>